

# Type State

CS242

Lecture 10

# APIs

Consider a typical use of a file API:

```
File f = open("MyFile");
```

```
Char c = f.read();
```

```
f.close()
```

# APIs

File operations mutate state and can change what operations are legal

```
File f = open("MyFile"); // opening the file enables reading
```

```
Char c = f.read();
```

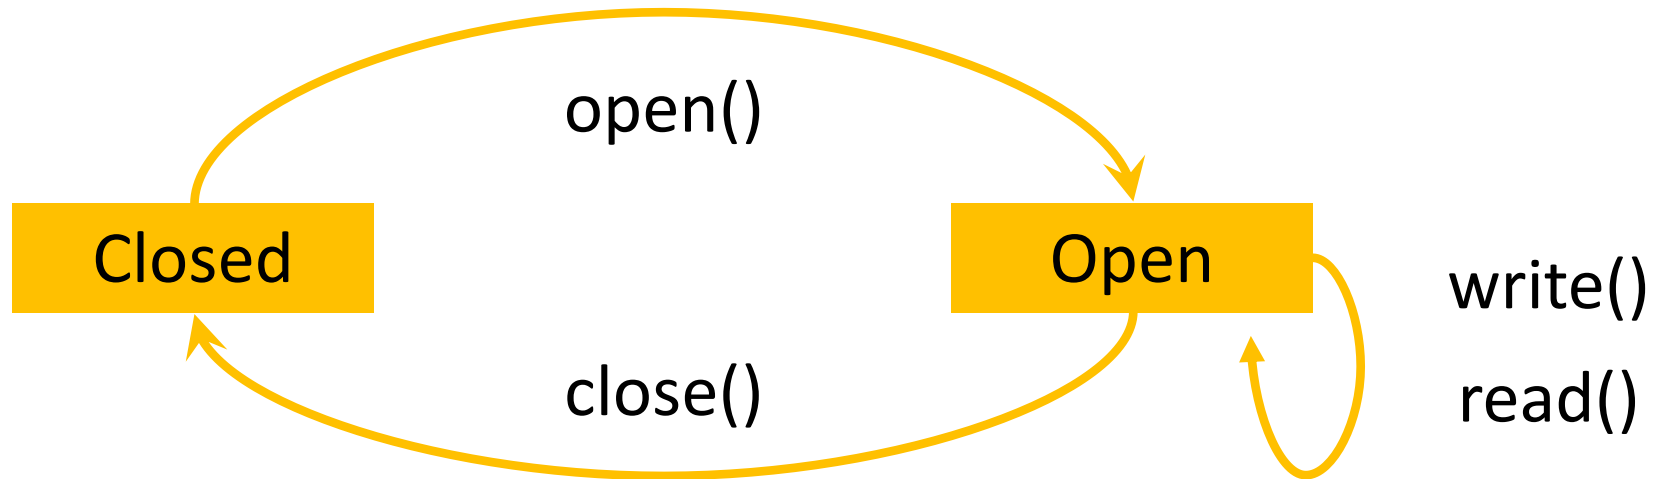
```
f.close() // but once the file is closed, reading is no  
// legal even though f is still available
```

# Stateful APIs

- Many APIs are stateful in this way
- Correct usage depends on the state of the system
  - Which changes over time as API calls are made
- States and state changes are commonly modeled using finite state machines

# Files

- Files can be open or closed
- A closed file can only be opened
- An open file can be read, written or closed

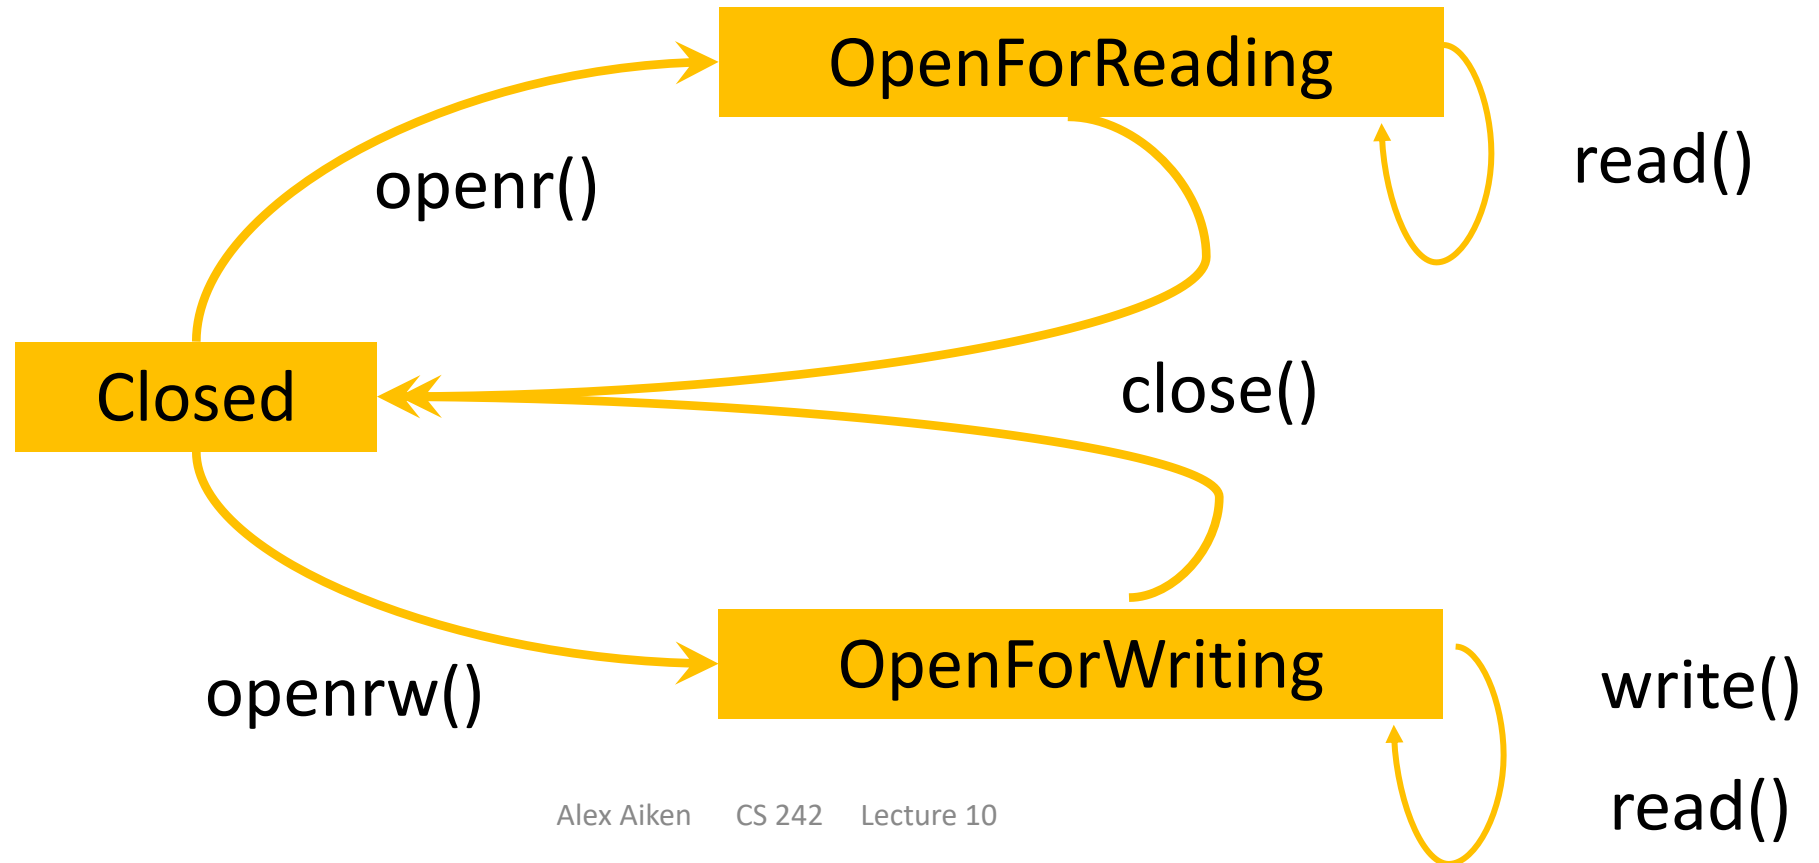


# Discussion

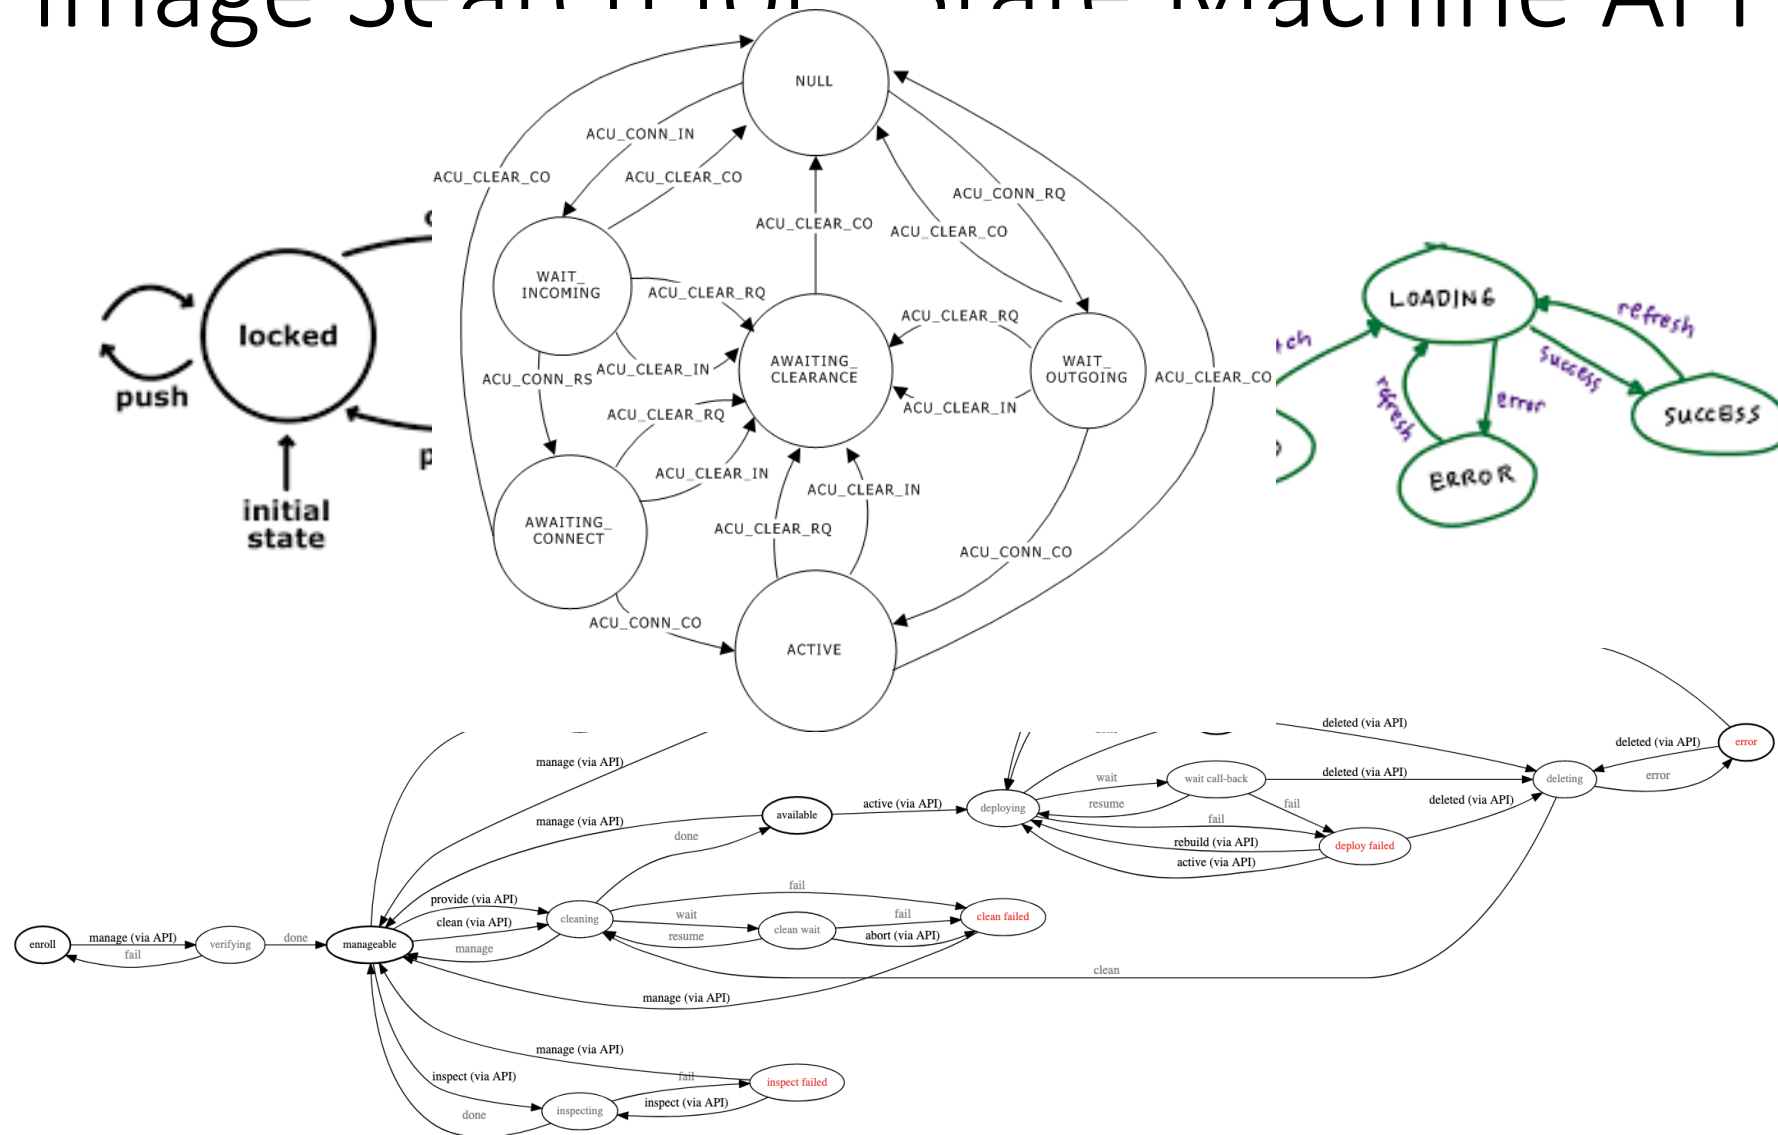
- Many similar examples of this two state machine in real APIs:
- File open/close
- Lock acquire/release
- Message encrypted/decrypted
- User logged in/logged out
- Resource initialized/uninitialized

# More Complex Machines

- File example can be refined, revealing more states specific to reading/writing



# Image Search for “State Machine API”





# Enforcing APIs

- We would like to enforce that API usage is correct according to a state machine specification
- Guarantees
  - We start in a correct state
  - The only calls that are made are valid in the current state
  - Can also enforce that we end in a correct state
    - e.g., Don't leave a resource locked forever

# Type State

- Idea: Use the type system
- Each state machine state corresponds to a distinct type
- Method invocations cause transitions in the state diagram
  - `open(): ClosedFile -> OpenFile`
  - `close(): OpenFile -> ClosedFile`
  - `read(): OpenFile -> OpenFile`

# A Possible Implementation

```
Class FileRecord {  
    File f;  
    Str name; }  

```

```
Class ClosedFile {  
    FileRecord r;  
    new ClosedFile(r: FileRecord);  
    open(): OpenFile; }  

```

```
Class OpenFile {  
    FileRecord r;  
    new OpenFile(name: Str);  
    read(): Char;  
    close(): ClosedFile }  

```

# Example

```
f = new OpenFile("foo")    // f is of type OpenFile  
c = f.read()  
x = f.close()              // x is of type ClosedFile
```

// but ...

```
d = f.read()  // f has been closed!
```

# Discussion

- The `OpenFile` and `ClosedFile` types allow us to express the transition from an open file to a closed file in the type system
  - And even back again
- The types only permit valid operations for the state of the file
- But ...
  - Nothing gets rid of the `OpenFile` objects when the file is closed
  - We could null out the `FileRecord`, but the `OpenFile` object would still exist

# What Do We Want?

```
f = new OpenFile("foo")
```

```
c = f.read()
```

```
x = f.close()
```

```
// ideally at this point the f object would be deallocated
```

```
// and it would also be impossible to invoke operations on f
```

```
// So...
```

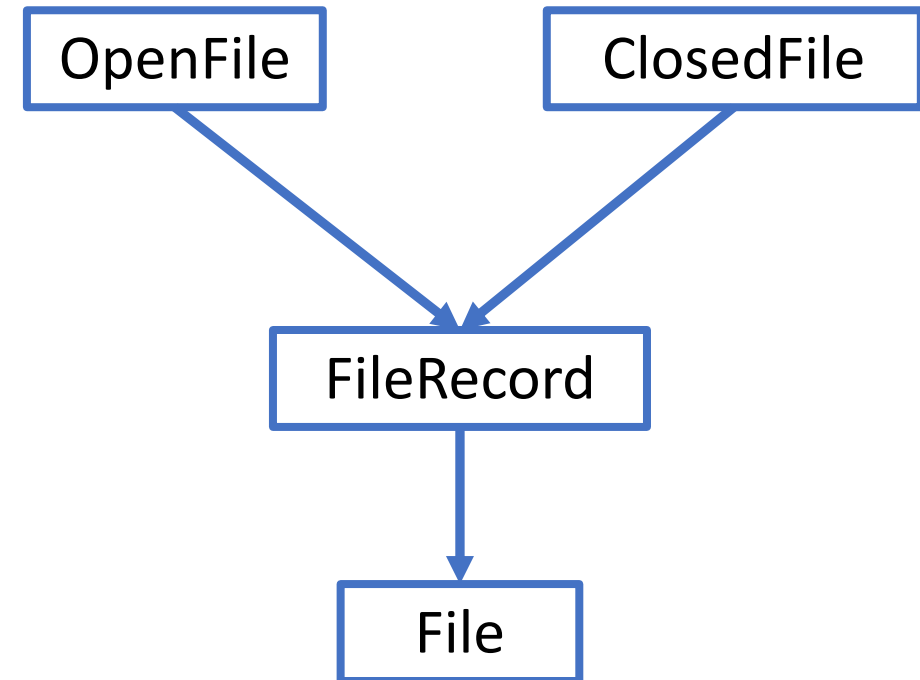
```
d = f.read() // this should be a type error!
```

# The Real Problem

- Types are not the problem
  - We can express the state transitions as functions that map one type to another
- Keeping track of object references is the problem
  - In particular, tracking *aliasing* between names

# Aliasing

- OpenFile and ClosedFile methods both modify the state of the File object in the FileRecord
- An example of *aliasing*
  - Having multiple names (ways to access) the same value
- Either object can change the state of the file without the knowledge of the other





# Aliasing Control

# A Classic Example

```
copy(char *x, char *y) {  
    ...  
}
```

But what about `copy(a,a)`?

# A Classic Example

```
copy(restrict char *x, restrict char *y) {  
    ...  
}
```

Semantics: In C, a restricted pointer cannot be aliased to any other pointer in scope.

# A Point of View

- Aliasing is bad
- State can be modified through one name and those changes are visible through a different name
  - Leads to subtle and difficult bugs
- But aliasing is very common in real programs
  - Impossible to avoid
  - E.g., references passed as arguments to functions
  - Object-oriented code is particularly prone to generating aliasing

# Idea #1

- Maybe aliasing is not the problem ...
- Problems arise only when aliasing is combined with mutation
- So, disallow mutation!
  - The pure functional programming viewpoint

# Could This Really Work?

- People have studied pure functional languages for decades
  - No mutation, whenever a data structure is changed a copy is made
- A surprising number of standard algorithms are just as efficient without mutation of state
  - Sometimes just amortized bounds, but that is still quite good!
- But there are some operations that seem to fundamentally require mutation to be efficient
  - Update in place of an array is  $O(1)$
  - The best known functional update is  $O(\log N)$  in the size of the array

# A Practical Approach

- Split the world into mutable and immutable values
- Rust
  - `let x = 5`      `// immutable`
  - `let mut x = 5` `// mutable`
  - `x = 3`      `// only allowed if x is mutable`
- ML
  - `let x = 5`      `// immutable`
  - `let x = ref 5` `// mutable`
  - `x := 3`

# Separating Mutable & Immutable

- Not entirely a new idea
  - E.g., `const` in C
- Gaining in popularity
  - More languages are making this distinction
  - With immutability being the default
- Now accepted as a good idea
  - Limit the possibility of mutation to places it is really needed
  - Make these points obvious in the syntax & types



# Idea #2

- Control aliasing in the type system
  - Track it, restrict it, or even disallow it
- Ownership types
  - Track aliases using types
- There is a large literature on ownership types
  - Some quite elaborate ...

# Ownership in Rust

- Rust has a simple ownership model
- There is always a single *owner* variable of every object
  - Owning = responsible for the resources of the object
- Implications
  - An object with no owner is deallocated
    - When an owner goes out of scope, the owned object is deallocated
  - Copies transfer ownership
    - $x = y$  removes ownership from  $y$  and transfers it to  $x$
    - $y$  can no longer be used after the assignment

# Ownership Example

```
fn main(){  
    let v = vec[1,2,3];    // v owns the vector  
    let v2 = v;            // moves ownership to v2  
    // let v3 = v[1] compile-time error!  
    display(v2);          // ownership is moved to display  
    // println(v2); compile-time error!  
}  
  
fn display(v:Vec<i32>){  
    println(v);  
    // v goes out of scope here and the vector is deallocated  
}
```

# Another Ownership Example

```
A() {  
  X = new Foo(); // X is the owner  
  Y = Bar(X)      // ownership is transferred to the argument of Bar  
                  // and then back to Y  
  // Y goes out of scope and the vector is deallocated  
}
```

```
Bar(Z) {  
  return Z // ownership is transferred back to the caller  
}
```

# Lifetimes

- Rust reasons about aliasing/ownership by using *lifetimes*
- The lifetime of a variable is the span between
  - The definition (first use)
  - The last use
- Roughly: Lifetimes of aliases cannot overlap
  - Thus it is important to minimize lifetimes
  - We will refine this rule shortly

# Lifetimes

```
fn main(){
```

```
    let v = vec[1,2,3]; // vector v owns the object
```

```
    let v2 = v;          // moves ownership to v2
```

```
    // let v3 = v[1] compile-time error!
```

```
    display(v2);         // ownership is moved to display
```

```
    // println(v2); compile-time error!
```

```
}
```

```
fn display(v:Vec<i32>){
```

```
    println(v);
```

```
    // v goes out of scope here and the vector is deallocated
```

```
}
```

# Lifetimes: A Compile Time Error

```
fn main(){  
    let v = vec[1,2,3]; // vector v owns the object  
    let v2 = v;         // moves ownership to v2  
    let v3 = v[1]       // compile-time error!  
    display(v2);        // ownership is moved to display  
    // println(v2); compile-time error!  
}
```

```
fn display(v:Vec<i32>){  
    println(v);  
    // v goes out of scope here and the vector is deallocated  
}
```

# Lifetimes: A Fix

```
fn main(){  
    let v = vec[1,2,3]; // vector v owns the object 2  
    let v3 = v[1]       // now this works ...  
    let v2 = v;         // moves ownership to v2  
    display(v2);        // ownership is moved to display  
    // println(v2); compile-time error!  
}
```

```
fn display(v:Vec<i32>){  
    println(v);  
    // v goes out of scope here and the vector is deallocated  
}
```



# Aliasing Control in Rust

- Having only a single owner is painful in many situations
  - Can never have another name for an object or even a piece of an object
  - E.g., makes it impossible to name a subarray
  - And we often don't need to take ownership anyway
- Rust allows the creation of explicit aliases
  - called *references* or *borrows*
- There are two kinds of references:
  - mutable
  - immutable

# Example: Immutable Reference

```
A() {  
  X = new Foo(); // X is the owner  
  Y = &X         // An immutable reference to X; X is still the owner.  
  Bar(Y)         // pass an immutable reference to Bar  
}
```

```
Bar(&Z) {  
  ... = .. Z ... // can read from Z in Bar as many times as we like  
  // Global.x = Z but if we try to store Z somewhere that outlives Bar we get an error  
}
```

# Example: Immutable Reference

```
A() {  
  X = new Foo(); // X is the owner  
  Y = &X         // An immutable reference to X; X is still the owner.  
  Bar(Y, Y)      // pass two immutable references to Bar  
}
```

```
Bar(&A, &B) {  
  ... = .. A ... // can read from A and B in Bar as many times as we like  
  ... = ... B ...  
  
}
```

# Example: Mutable Reference

```
A() {  
  X = new Foo();    // X is the owner  
  Y = &mut X;       // Y is a mutable reference to X  
  Bar(Y)            // pass a mutable reference to Bar  
}
```

```
Bar(&mut Z) {  
  Z.f = ... // can mutate Z  
}
```

# Example: Mutable Reference

```
A() {  
  X = new Foo();    // X is the owner  
  Y = &mut X;       // Y is a mutable reference to X  
  Bar(Y, Y)         // Error: Cannot have two mutable references to X  
}
```

```
Bar(&mut A, &mut B) { // since A and B are mutable, they cannot alias  
  A.f = ... // can mutate A  
  B.f = ... // can mutate B  
}
```

# Reference Rules

- A reference cannot outlive its referent
  - No dangling references ...
- A mutable reference cannot be aliased
- Meaning:
  - There can be one mutable reference to an object in scope
  - There can be any number of immutable references
  - The owner's lifetime must contain the lifetime of all references

# Example: Immutable Reference

```
A() {
```

```
  X = new Foo(); // X is the owner
```

```
  Y = &X          // An immutable reference to X; X is still the owner.
```

```
  Bar(Y)          // pass an immutable reference to Bar; the reference's lifetime is the lifetime of Bar
```

```
}
```

```
Bar(&Z) {
```

```
  ... = .. Z ...    // can read from Z in Bar as many times as we like
```

```
  // Global.x = Z    but if we try to store Z somewhere that outlives Bar we'll get an error
```

```
}
```

# Back to Type State

- Recall type state allowed us to express state changes in the type system:

```
// f: OpenFile
```

```
x = f.close() // x is of type ClosedFile
```

- Ownership types can be used to make objects corresponding to the previous API state inaccessible



# The Type State Example

```
f = new OpenFile("foo")    // f is the owner
```

```
c = f.read()
```

```
x = f.close()  // x is of type ClosedFile, transfers file ownership f -> x
```

```
// then ...
```

```
d = f.read()  // generates a type error: f is not the owner
```

# Type State Abstraction

- Objects of a given state support only the methods available in that state
- Methods that change the internal API state
  - Return a new object of the appropriate type
  - Transfer ownership to the new object so that the references to the previous state become inaccessible
- Thus only objects of the correct type are ever accessible by the program

# Discussion

- Ownership rules are very restrictive
  - Program must be *linear* in owned objects
  - Exactly one owner at all times
  - Implies no copies can be made
- Three techniques help:
  - Using immutable data wherever possible
  - Deep copies are OK (*cloning*)
  - Borrowing creates a reference that can be used
    - Does not transfer ownership
    - Implies a borrowed reference cannot deallocate an object
    - The owner cannot deallocate an object until all borrowed references are returned
    - Borrowed references have a different syntax and type

# Ownership in Practice

- Ownership has been studied for > 20 years
- Rust is the first full language to support ownership types
  - The major new feature
- Experience is that Rust's ownership system helps
  - Enables manually managed memory without the bugs
  - Makes it possible to write efficient and correct code
  - Ownership types are the key
  - Which is not to say ownership is always easy to use ... it's not!